

431 Class 17

Thomas E. Love, Ph.D.

2022-11-01

Today's Agenda

- Power and Sample Size: An Introduction
 - When Comparing Two Means
 - When Comparing Two Proportions

Today's Packages

```
1 library(pwr)
2 library(tidyverse)
3
4 theme_set(theme_bw())
```

How Big A Sample Size Do I need?

1. What is the budget?
2. What are you trying to compare?
3. What is the study design?
4. How big an effect size do you expect (hope) to see?
5. What was that budget again?
6. OK, tell me the maximum allowable rates of Type I and Type II error that you want to control for. Or, if you like, tell me the confidence level and power you want to have.
7. What sort of statistical inference do you want to plan for?

Errors in Hypothesis Testing

In testing hypotheses, there are two potential decisions and each one brings with it the possibility that a mistake has been made.

	H_0 is true	H_A is true
Test rejects H_0	Type I error (False positive)	Correct
Test retains H_0	Correct	Type II error (False negative)

Errors in Hypothesis Testing

- A Type I error can only be made if (H_0) is actually true.
- A Type II error can only be made if (H_A) is actually true.

	(H_A) is true	(H_0) is true
Test rejects (H_0)	Correct	Type I error (False positive)
Test retains (H_0)	Type II error (False negative)	Correct

Specifying Error Probabilities (Type I)

If we say we are using 90% confidence, this means:

- we have a 10% significance level
- α , the probability of Type I error, is set to 0.10
- In general, confidence level = $100(1 - \alpha)$.
- The probability of correctly retaining H_0 is designed to be 0.90.

Specifying Error Probabilities (Type II)

Trading off significance (α) and power ($1 - \beta$).

In many sample size decisions,

- we find that people set α , the tolerable rate of Type I error, to be 0.05.
- they then often try to set the sample size and other parameters so that the power ($1 - \beta$) is at least 0.80.

Relative Costs of Errors

We'll advocate for thinking hard about the relative costs of Type I and Type II errors.

- The underlying framework that assumes a power of 80% with a significance level of 5% is sufficient for most studies is pretty silly.

Power and Sample Size Calculations

A power calculation is likely the most common element of an scientific grant proposal on which a statistician is consulted.

- The tests that have power calculations worked out in intensive detail using R are mostly those with more substantial assumptions.
 - t tests that assume population normality, common population variance and balanced designs in the independent samples setting
 - paired t tests that assume population normality

Power and Sample Size Calculations

- These power calculations are also usually based on tests rather than confidence intervals. Simulation is your friend here.
- This process of doing power and related calculations is far more of an art than a science.

Power and Sample Size: Comparing Means

Paired vs. Independent Samples

If you can afford to obtain $n = 400$ observations to compare means under exposure A to means under exposure B, and you could either:

1. select a random sample from the population of interest containing 400 people and then randomly assign 200 people to receive exposure A and the remaining 200 people to receive exposure B (thus doing an independent samples study), or
2. select a random sample from the population of interest containing 200 people and then randomly assign 100 of them to get exposure A first, and then, a little later, when the effects have worn off, to then receive exposure B, while the other 100 people are assigned to receive B first, then A (thus doing a paired samples study)

Assuming the effect size is unchanged, which seems as though it would be the more powerful study design?

Power of an Independent Samples t test

```
1 power.t.test(n = 200, delta = 0.25, sd = 1,  
2             sig.level = 0.10)
```

Two-sample t test power calculation

```
      n = 200  
    delta = 0.25  
      sd = 1  
sig.level = 0.1  
  power = 0.8025858  
alternative = two.sided
```

NOTE: n is number in *each* group

Power of an Paired Samples t test

```
1 power.t.test(n = 200, delta = 0.25, sd = 1,  
2             sig.level = 0.10, type = "paired")
```

Paired t test power calculation

```
      n = 200  
    delta = 0.25  
      sd = 1  
sig.level = 0.1  
  power = 0.9698521  
alternative = two.sided
```

NOTE: n is number of *pairs*, sd is std.dev. of *differences* within pairs

What sample size do we need?

How many pairs of observations would we need to maintain 80% power?

```
1 power.t.test(delta = 0.25, sd = 1, sig.level = 0.10,  
2               power = 0.80, type = "paired")
```

Paired t test power calculation

```
      n = 100.2877  
delta = 0.25  
    sd = 1  
sig.level = 0.1  
  power = 0.8  
alternative = two.sided
```

NOTE: n is number of *pairs*, sd is std.dev. of *differences* within pairs

- Note that we'd need 101 pairs of measurements.

Changing assumptions?

In our independent-samples test, we chose

- `n = 200` (per group)
- `delta = 0.25` (the minimum clinically important difference in means that we want to detect)
- `sd = 1` (assumed population standard deviation in each group)
- `sig.level = 0.10` (since we want 90% confidence)

Which direction will power move in?

Original Setup yielded power = 0.802

- If we change n from 200 to 400, leaving everything else untouched, do you think the power will increase or decrease?
- If we change n from 200 to 400, power = 0.970
- What if we change n from 200 to 100?
- If we change n from 200 to 100, power = 0.546

Changing other parameters

Original power = 0.802. Which changes **increase** the power?

New Setup	Resulting Power
a. Δ from 0.25 to 0.5	Higher or Lower than 0.802?
b. Δ from 0.25 to 0.1	?
c. sd from 1 to 2	?
d. sd from 1 to 0.5	?
e. α from 0.1 to 0.05	?

Changes a, d, f increase power

Original Setup yielded power = 0.802

Change	Resulting power
a. Δ from 0.25 to 0.5	0.9996
b. Δ from 0.25 to 0.1	0.259
c. sd from 1 to 2	0.345
d. sd from 1 to 0.5	0.9996
e. α from 0.1 to 0.05	0.703
f. α from 0.1 to 0.2	0.888

What if you have an unbalanced design?

The most efficient design for an independent samples comparison will be balanced.

- What if we used our original setup for δ , sd and α , (which with $n = 200$ in each group, yielded power = 0.802) but instead we placed
 - 150 subjects into one exposure group, and
 - planned to recruit some number X larger than 150 into the other.

Unbalanced Design?

- How many people would we have to recruit into the second exposure group to yield the same power as our original 200 in each group result?

`pwr.t2n.test` from the `pwr` package

- Note the use here of $d = \frac{\Delta}{sd}$.

```
1 pwr.t2n.test(n1 = 150, d = 0.25/1,  
2             sig.level = 0.10, power = 0.802)
```

t test power calculation

```
      n1 = 150  
      n2 = 298.1132  
      d = 0.25  
sig.level = 0.1  
  power = 0.802  
alternative = two.sided
```

So we can either have 200 and 200, or we can have 150 and 299 to maintain the same power.

Assessing Unbalanced Designs

The power is always stronger for a balanced design than for an unbalanced design with the same overall sample size.

See chapter 22 of the Course Notes for additional examples using the `pwr.t2n.test()` function within the `pwr` package.

One-Sided or Two-Sided

Note that I used a two-sided test to establish my power calculation - in general, this is the most conservative and defensible approach for any such calculation, unless there is a strong and specific reason to use a one-sided approach in building a power calculation, don't.

Power and Sample Size: Comparing Proportions

Designing a New TB Study

(PI): OK. That's a nice pilot. We saw $(p_{\text{nonshare}}) = 0.18$ and $(p_{\text{share}}) = 0.26$ after your augmentation. Help me design a new study using a two-sided test with $(\alpha = 0.05)$.

- This time, let's have as many needle-sharers as non-sharers.
- We should have 90% power to detect a difference almost as large as what we saw in the pilot, or larger, so a difference of 6 percentage points.

How `power.prop.test` works

We specify 4 of the following 5 elements of the comparison, and R calculates the fifth.

- The sample size (interpreted as the # in each group, so half the total sample size)
- The true probability in group 1
- The true probability in group 2
- The significance level (α)
- The power ($1 - \beta$)

Requires you to work with balanced designs.

Using `power.prop.test`

To find the sample size for a two-sample comparison of proportions using a balanced design:

- we will use a two-sided test, with $\alpha = .05$, and power = .90,
- we estimate that non-sharers have probability .18 of positive tests,
- and we will try to detect a difference between this group and the needle sharers, who we estimate will have a probability of .24

Any guess as to needed sample size?

Finding the required sample size in R

```
1 power.prop.test(p1 = .18, p2 = .24, alternative = "two.sided",  
2               sig.level = 0.05, power = 0.90)
```

Two-sample comparison of proportions power calculation

```
      n = 966.3554  
    p1 = 0.18  
    p2 = 0.24  
sig.level = 0.05  
  power = 0.9  
alternative = two.sided
```

NOTE: n is number in *each* group

Sample Size Required

So, we'd need at least 967 non-sharing subjects, and 967 more who share needles to accomplish the aims of the study, or a total of 1934 subjects.

Another Scenario

Suppose we can get 400 sharing and 400 non-sharing subjects. With what power could we detect 0.18 in one group and 0.26 in the other, in a *one-sided* $(\alpha) = .10$ test?

```
1 power.prop.test(n=400, p1=.18, p2=.26, sig.level = 0.10,
2                 alternative="one.sided")
```

Two-sample comparison of proportions power calculation

```
      n = 400
      p1 = 0.18
      p2 = 0.26
sig.level = 0.1
  power = 0.9273602
alternative = one.sided
```

NOTE: n is number in *each* group

Using the `pwr` package to assess sample size for Unbalanced Designs

The `pwr.2p2n.test` function in the `pwr` package can help assess the power of a test to determine a particular effect size using an unbalanced design, where (n_1) is not equal to (n_2) .

As before, we specify four of the following five elements of the comparison, and R calculates the fifth.

Now the five elements are...

- **n1** = The sample size in group 1
- **n2** = The sample size in group 2
- **sig.level** = The significance level (α)
- **power** = The power ($1 - \beta$)
- **h** = the effect size, which can be calculated separately in R based on the two proportions being compared: p_1 and p_2 .

Calculating the Effect Size h

To calculate the effect size for a given set of proportions, use `ES.h(p1, p2)` which is available in the `pwr` package.

For instance, comparing .18 to .25, we have the following effect size.

```
1 ES.h(p1 = .18, p2 = .25)
```

```
[1] -0.1708995
```

Using `pwr.2p2n.test` in R

Suppose we can have 700 samples in group 1 (the not sharing group) but only 400 in group 2 (the group of users who share needles).

How much power would we have to detect the distinction between $p_1 = .18$, $p_2 = .25$ with a 5% significance level in a two-sided test?

R Command to find the resulting power

```
1 pwr::pwr.2p2n.test(h = ES.h(p1 = .18, p2 = .25), n1 = 700, n2 = 400,  
2                      sig.level = 0.05)
```

difference of proportion power calculation for binomial distribution
(arcsine transformation)

```
      h = 0.1708995  
     n1 = 700  
     n2 = 400  
sig.level = 0.05  
   power = 0.7783562  
alternative = two.sided
```

NOTE: different sample sizes

Comparison to Balanced Design

How does this compare to the results with a balanced design using 1100 drug users in total, i.e. 550 per group?

```
1 pwr::pwr.2p2n.test(h = ES.h(p1 = .18, p2 = .25), n1 = 550, n2 = 550,
2                      sig.level = 0.05)
```

difference of proportion power calculation for binomial distribution
(arcsine transformation)

```
      h = 0.1708995
     n1 = 550
     n2 = 550
sig.level = 0.05
   power = 0.8089644
alternative = two.sided
```

NOTE: different sample sizes

We could instead have used...

```
1 power.prop.test(p1 = .18, p2 = .25, sig.level = 0.05, n = 550)
```

Two-sample comparison of proportions power calculation

```
      n = 550
     p1 = 0.18
     p2 = 0.25
sig.level = 0.05
  power = 0.8075197
alternative = two.sided
```

NOTE: n is number in *each* group

Not the Same?

Each approach uses approximations, and slightly different ones, so it's not surprising that the answers are similar, but not identical.

What haven't I included here?

1. Some people will drop out.
2. What am I going to do about missing data?
3. What if I want to do my comparison while adjusting for covariates?

More examples (again not touching on these last three issues much) in Chapter 27 of the Course Notes.

On Power / Sample Size Decisions and the March of Science

Power Calculations

The most common scenario was to identify the desired significance level α and desired power $(1 - \beta)$ and the details of the plan in terms of what comparison is to be made, and how will the data be collected to support that comparison.

This will then permit the calculation of a minimum necessary sample size to achieve these desires.

$\alpha = 0.05$ and $\beta = 0.2$ are the most common selections.

- Neither 95% confidence nor 80% power is a magical choice.
- Anything below 80% power will be hard to justify in real work.

A useful metaphor?

Sometimes I like to think of science as a march towards a destination.

Actually, I suppose it's an infinitely long march towards an ever-receding destination, but let's leave the philosophy out of it for a moment.

Suppose, for example, that we're trying to make a meaningful change in the world, perhaps to treat an infection.

What we're trying to do is related to where we are in the March of Science.

Early vs. Late in the March of Science

In **early** work, we're focused more on discovery than making final decisions.

- We don't have a lot of past experience, so we bring little relevant data to the table.
- We're (often) most concerned about discovering new possibilities, and we don't have a very clear sense of where to go next.
- We're (often) less concerned about false starts than we are about missed opportunities.

In **late** work, we're focused more on making a decision about how to treat.

- We have a fair amount of relevant history to draw on, sometimes quite detailed.
- We're more concerned about testing the limits of our current knowledge than we are about missing opportunities to consider a new pathway.
- We're often concerned about doing harm if we implement the strategy that looks most promising.

Power and “significance”?

If we treat our sample size and study design as fixed strategies, then there is a tradeoff between:

- reducing α , the rate of Type I error (increasing our confidence) and
- reducing β , the rate of Type II error (increasing our power)

Suppose we are testing a new treatment for some condition.

- A Type I error means we conclude this treatment is helpful, when it actually isn't.
- A Type II error means we conclude this treatment is not helpful, when it actually is.

Early Work: Power and Sample Size

In early work, we are searching for treatments of promise, and our initial study will inevitably not be the last word on the subject, but rather will be followed up by confirmatory studies. In such a setting, it is often the case that:

- We're not so concerned about getting results that cause us to continue to explore a treatment that doesn't actually do what we need it to do.
- We're really concerned about ruling out a treatment that is promising before we should.

This implies we should prioritize reducing Type II error rates (we want more power to detect small but real effects, even if this means we will occasionally identify something as promising when it isn't.)

This means setting lower confidence levels and higher power levels, potentially, than the standard 95% confidence and 80% power.

Late Work: Power and Sample Size

In late work, we have already identified promising treatments, and we are trying to confirm those results. The current study may actually be the last word on the subject, and we want to be sure we do no harm.

- We're very concerned about getting results that cause us to continue to explore a treatment that doesn't actually do what we need it to do.
- We're less concerned about ruling out a treatment that is promising but doesn't actually work.

This implies we should prioritize reducing Type I error rates (we want greater confidence, even at the expense of power, that the effect we claim based on past data holds up.)

This kind of confirmatory work is usually well suited to studies set up with higher confidence (perhaps 95% or 99% or more) and lower power (80% is the minimum I would recommend) against reasonable alternatives.

Conclusions

1. If you're early in the March of Science (perhaps just one pilot study has been done) then I would emphasize Type II error (power) more than usual, perhaps pushing required power to 90%, at least for a reasonably substantial "minimum scientifically important difference" δ .
2. If you're late in the March of Science (perhaps confirming the results of multiple prior studies) then I would be happier with 80% power and higher levels of confidence.
3. If you're in the middle, trading off Type I and Type II error is worth some thought, but I'd never recommend being under 80% power for an effect that matters.
4. If it's feasible to run a study large enough to have strong performance on both α and β , that's obviously ideal. Typically that doesn't happen in early work.

Session Information

```
1 sessionInfo()
```

```
R version 4.2.1 (2022-06-23 ucrt)  
Platform: x86_64-w64-mingw32/x64 (64-bit)  
Running under: Windows 10 x64 (build 22000)
```

```
Matrix products: default
```

```
locale:
```

```
[1] LC_COLLATE=English_United States.utf8  
[2] LC_CTYPE=English_United States.utf8  
[3] LC_MONETARY=English_United States.utf8  
[4] LC_NUMERIC=C  
[5] LC_TIME=English_United States.utf8
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```